

Week 2 Data Exploration

```
In [1]: # This code appears in every demonstration Notebook.
# By default, when you run each cell, only the last output of the codes will show.
# This code makes all outputs of a cell show.
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
```

We will explore the user and order data from JD.com

1. We import the necessary packages.

```
In [2]: import pandas as pd
```

2. We read in the datasets.

```
In [3]: users = pd.read_csv('JD_user_data.csv')
orders = pd.read_csv('JD_order_data.csv')
```

3. Take a look at the data

```
In [4]: users.head()
orders.head()
```

```
Out[4]:
```

| | user_ID | user_level | first_order_month | plus | gender | age | marital_status | education |
|---|------------|------------|-------------------|------|--------|-------|----------------|-----------|
| 0 | 000089d6a6 | 1 | 2017-08 | 0 | F | 26-35 | S | |
| 1 | 0000babd1f | 1 | 2018-03 | 0 | U | U | U | - |
| 2 | 0000bc018b | 3 | 2016-06 | 0 | F | >=56 | M | |
| 3 | 0000d0e5ab | 3 | 2014-06 | 0 | M | 26-35 | M | |
| 4 | 0000dce472 | 3 | 2012-08 | 1 | U | U | U | - |

Out[4]:

| | order_ID | user_ID | sku_ID | order_date | order_time | quantity | type | promise |
|---|------------|------------|------------|------------|-----------------------|----------|------|---------|
| 0 | d0cf5cc6db | 0abe9ef2ce | 581d5b54c1 | 2018-03-01 | 2018-03-01 17:14:25.0 | 1 | 2 | - |
| 1 | 7444318d01 | 33a9e56257 | 067b673f2b | 2018-03-01 | 2018-03-01 11:10:40.0 | 1 | 1 | 2 |
| 2 | f973b01694 | 4ea3cf408f | 623d0a582a | 2018-03-01 | 2018-03-01 09:13:26.0 | 1 | 1 | 2 |
| 3 | 8c1cec8d4b | b87cb736cb | fc5289b139 | 2018-03-01 | 2018-03-01 21:29:50.0 | 1 | 1 | 2 |
| 4 | d43a33c38a | 4829223b6f | 623d0a582a | 2018-03-01 | 2018-03-01 19:13:37.0 | 1 | 1 | 1 |

In [5]: `users.head(10) # The number argument specifies the number of rows to show`

Out[5]:

| | user_ID | user_level | first_order_month | plus | gender | age | marital_status | education |
|---|------------|------------|-------------------|------|--------|-------|----------------|-----------|
| 0 | 000089d6a6 | 1 | 2017-08 | 0 | F | 26-35 | S | |
| 1 | 0000babd1f | 1 | 2018-03 | 0 | U | U | U | - |
| 2 | 0000bc018b | 3 | 2016-06 | 0 | F | >=56 | M | |
| 3 | 0000d0e5ab | 3 | 2014-06 | 0 | M | 26-35 | M | |
| 4 | 0000dce472 | 3 | 2012-08 | 1 | U | U | U | - |
| 5 | 0000f81d1b | 1 | 2018-02 | 0 | F | 26-35 | M | |
| 6 | 00012bb423 | 4 | 2008-11 | 1 | F | 26-35 | M | |
| 7 | 00015ff032 | 3 | 2015-06 | 1 | M | 26-35 | M | |
| 8 | 0001aa7059 | 4 | 2014-06 | 0 | F | 36-45 | M | |
| 9 | 0001bbdc89 | 2 | 2017-12 | 0 | F | 16-25 | S | |



```
In [6]: users.columns  
# Displays the variables of the dataframe
```


```
Out[6]: Index(['user_ID', 'user_level', 'first_order_month', 'plus', 'gender', 'age',  
             'marital_status', 'education', 'city_level', 'purchase_power'],  
            dtype='object')
```

```
In [7]: # Lets explore the gender variable in users  
users['gender'].value_counts()  
# value_counts() gives the frequency distribution
```

```
Out[7]: gender  
F      292897  
M      107084  
U       57317  
Name: count, dtype: int64
```

```
In [8]: # Make a bar chart for the frequency distribution  
users['gender'].value_counts().plot(kind = 'bar')
```


```
Out[8]: <Axes: xlabel='gender'>
```



```
In [9]: # Exercise: exploring age  
age_dis = users['age'].value_counts()
```

```
In [10]: users['age'].value_counts().plot(kind = 'bar')
```

```
Out[10]: <Axes: xlabel='age'>
```



```
In [11]: import matplotlib.pyplot as plt
```

```
#Importing the graph package matplotlib
```

```
In [12]: # Sample data
categories = ['26-35', '16-25', '36-45', 'U', '46-55', '>=56', '<=15']
values = [25000, 50000, 75000, 100000, 125000, 150000, 175000]

# Create a bar graph
plt.bar(categories, values, color='black')

# Add Labels and title
plt.xlabel('Categories')
plt.ylabel('Values')
plt.title('Age graph')


# Show the plot
plt.show()
```

```
Out[12]: <BarContainer object of 7 artists>
```

```
Out[12]: Text(0.5, 0, 'Categories')
```

```
Out[12]: Text(0, 0.5, 'Values')
```

```
Out[12]: Text(0.5, 1.0, 'Age graph')
```



In [13]: `age_dis.index`

Out[13]: `Index(['26-35', '16-25', '36-45', 'U', '46-55', '>=56', '<=15'], dtype='object', name='age')`

In [14]: `age_dis.values`

Out[14]: `array([183239, 103306, 81076, 56457, 18679, 14517, 24], dtype=int64)`

```
In [15]: # Sample data
custom_color = ['Red', 'Yellow', 'Purple', 'Orange', 'Blue', 'Green']
age = age_dis.index
values = age_dis.values


# Create a bar graph
plt.bar(categories, values, color= custom_color)

# Add labels and title
plt.xlabel('age')
plt.ylabel('Values')
plt.title('Age graph')

# Show the plot
plt.show()
```

Out[15]: `<BarContainer object of 7 artists>`

```
Out[15]: Text(0.5, 0, 'age')
Out[15]: Text(0, 0.5, 'Values')
Out[15]: Text(0.5, 1.0, 'Age graph')
```



5. Explore a single interval variable

```
In [16]: orders.columns
```

```
Out[16]: Index(['order_ID', 'user_ID', 'sku_ID', 'order_date', 'order_time', 'quantity',
               'type', 'promise', 'original_unit_price', 'final_unit_price',
               'direct_discount_per_unit', 'quantity_discount_per_unit',
               'bundle_discount_per_unit', 'coupon_discount_per_unit', 'gift_item',
               'dc_ori', 'dc_des'],
              dtype='object')
```

```
In [17]: orders['original_unit_price'].describe()
```

```
Out[17]: count    549989.000000
         mean      102.813542
         std       95.035563
         min        0.000000
         25%       59.000000
         50%       79.000000
         75%      139.000000
         max      12158.000000
         Name: original_unit_price, dtype: float64
```

```
In [18]: import numpy as np
```

```
In [19]: np.var(orders['original_unit_price'])
```

```
Out[19]: 9031.741770278562
```

```
In [20]: np.percentile(orders['original_unit_price'], 90) # Finding quantile or percentile
```

```
Out[20]: 240.0
```

```
In [21]: # Find the records with the maximum price
# The max price is an outlier
orders['original_unit_price'] == 12158
```

```
Out[21]: 0      False
1      False
2      False
3      False
4      False
...
549984  False
549985  False
549986  False
549987  False
549988  False
Name: original_unit_price, Length: 549989, dtype: bool
```


```
In [22]: orders[orders['original_unit_price'] == 12158]
```

```
Out[22]:
```

| | order_ID | user_ID | sku_ID | order_date | order_time | quantity | type | promi |
|--------------|------------|------------|------------|------------|-----------------------|----------|------|-------|
| 52720 | a63239c796 | b695af3c92 | 1904d943c0 | 2018-03-03 | 2018-03-03 14:56:01.0 | 1 | 2 | |

```
In [23]: # A histogram to explore the distribution of the interval variable
# original price
plt.hist(orders['original_unit_price'])
```

```
Out[23]: (array([5.49971e+05, 5.00000e+00, 2.00000e+00, 1.00000e+00, 2.00000e+00,
2.00000e+00, 2.00000e+00, 0.00000e+00, 1.00000e+00, 3.00000e+00]),
array([ 0. , 1215.8, 2431.6, 3647.4, 4863.2, 6079. , 7294.8,
8510.6, 9726.4, 10942.2, 12158. ]),
<BarContainer object of 10 artists>)
```



```
In [38]: orders = orders[orders['original_unit_price'] != 12158]
```

```
In [40]: orders['original_unit_price'].quantile(0.99)
```

```
Out[40]: 336.0
```

```
In [43]: orders['original_unit_price'].mean()
```

```
Out[43]: 102.7916229960169
```

```
In [42]: orders['original_unit_price'].std()
```

```
Out[42]: 93.63512178749312
```

```
In [31]: out_up = orders['original_unit_price'].mean() + 3 * orders['original_unit_price'].s
```

```
In [32]: print(out_up)
```

```
383.69698835849624
```


```
In [33]: #Delete the outliers using slicing  
normal_price = orders[orders['original_unit_price'] < 388]
```

```
In [34]: normal_price
```


Out[34]:

| | order_ID | user_ID | sku_ID | order_date | order_time | quantity | type | prc |
|---------------|------------|------------|------------|------------|--------------------------|----------|------|-----|
| 0 | d0cf5cc6db | 0abe9ef2ce | 581d5b54c1 | 2018-03-01 | 2018-03-01 17:14:25.0 | 1 | 2 | |
| 1 | 7444318d01 | 33a9e56257 | 067b673f2b | 2018-03-01 | 2018-03-01 11:10:40.0 | 1 | 1 | |
| 2 | f973b01694 | 4ea3cf408f | 623d0a582a | 2018-03-01 | 2018-03-01 09:13:26.0 | 1 | 1 | |
| 3 | 8c1cec8d4b | b87cb736cb | fc5289b139 | 2018-03-01 | 2018-03-01 21:29:50.0 | 1 | 1 | |
| 4 | d43a33c38a | 4829223b6f | 623d0a582a | 2018-03-01 | 2018-03-01 19:13:37.0 | 1 | 1 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 549984 | 3ad06b9fbe | a27b3ed4d4 | a9109972d1 | 2018-03-31 | 2018-03-31 01:22:47.0 | 1 | 2 | |
| 549985 | c9d77a7ed0 | 18f92434cd | 7f53769d3f | 2018-03-31 | 2018-03-31 08:55:57.0 | 1 | 1 | |
| 549986 | b9ad79338f | b5caf8a580 | 8dc4a01dec | 2018-03-31 | 2018-03-31 13:31:01.0 | 1 | 1 | |
| 549987 | be3a9414b1 | 20ba6655f3 | 2dd6b818ec | 2018-03-31 | 2018-03-31 12:51:18.0 | 1 | 2 | |
| 549988 | 02d31f05c9 | f260895cbe | 10d369ef96 | 2018-03-31 | 2018-03-31 18:21:16.0 | 1 | 2 | |

546883 rows × 17 columns



```
In [35]: orders.shape
#before slicing
```


Out[35]: (549988, 17)

```
In [44]: normal_price.shape
#after slicing
```

Out[44]: (546883, 17)

```
In [45]: plt.hist(normal_price['original_unit_price'])
```

```
Out[45]: (array([107222., 119175., 143434., 69727., 35921., 16537., 16975.,
        31718., 5087., 1087.]),
array([ 0., 38.5, 77., 115.5, 154., 192.5, 231., 269.5, 308.,
        346.5, 385. ]),
<BarContainer object of 10 artists>)
```



Exploration of two interval variables

```
In [ ]:
```

```
In [46]: normal_price.columns
```

```
Out[46]: Index(['order_ID', 'user_ID', 'sku_ID', 'order_date', 'order_time', 'quantity',
               'type', 'promise', 'original_unit_price', 'final_unit_price',
               'direct_discount_per_unit', 'quantity_discount_per_unit',
               'bundle_discount_per_unit', 'coupon_discount_per_unit', 'gift_item',
               'dc_ori', 'dc_des'],
              dtype='object')
```

```
In [51]: import seaborn as sns
```

```
In [52]: #Make a scatterplot between 'original_unit_price', 'final_unit_price',
plt.scatter(normal_price.original_unit_price, normal_price.final_unit_price)
sns.regplot(x = 'original_unit_price', y = 'final_unit_price', data = normal_price,
plt.title('Relationship between original_unit_price and final_unit_price'))
```

```
plt.xlabel('Original Price')
plt.ylabel('Final Price')
```

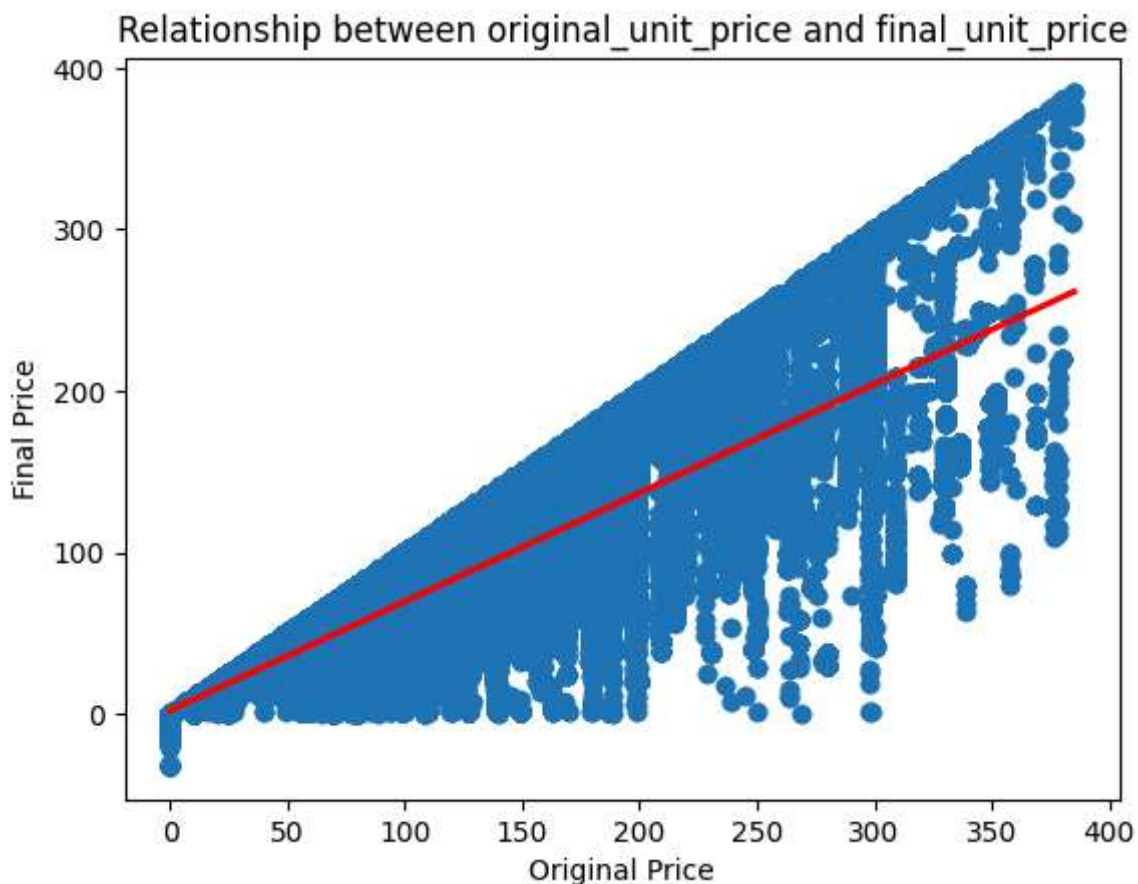
Out[52]: <matplotlib.collections.PathCollection at 0x209960d9b20>

Out[52]: <Axes: xlabel='original_unit_price', ylabel='final_unit_price'>

Out[52]: Text(0.5, 1.0, 'Relationship between original_unit_price and final_unit_price')

Out[52]: Text(0.5, 0, 'Original Price')

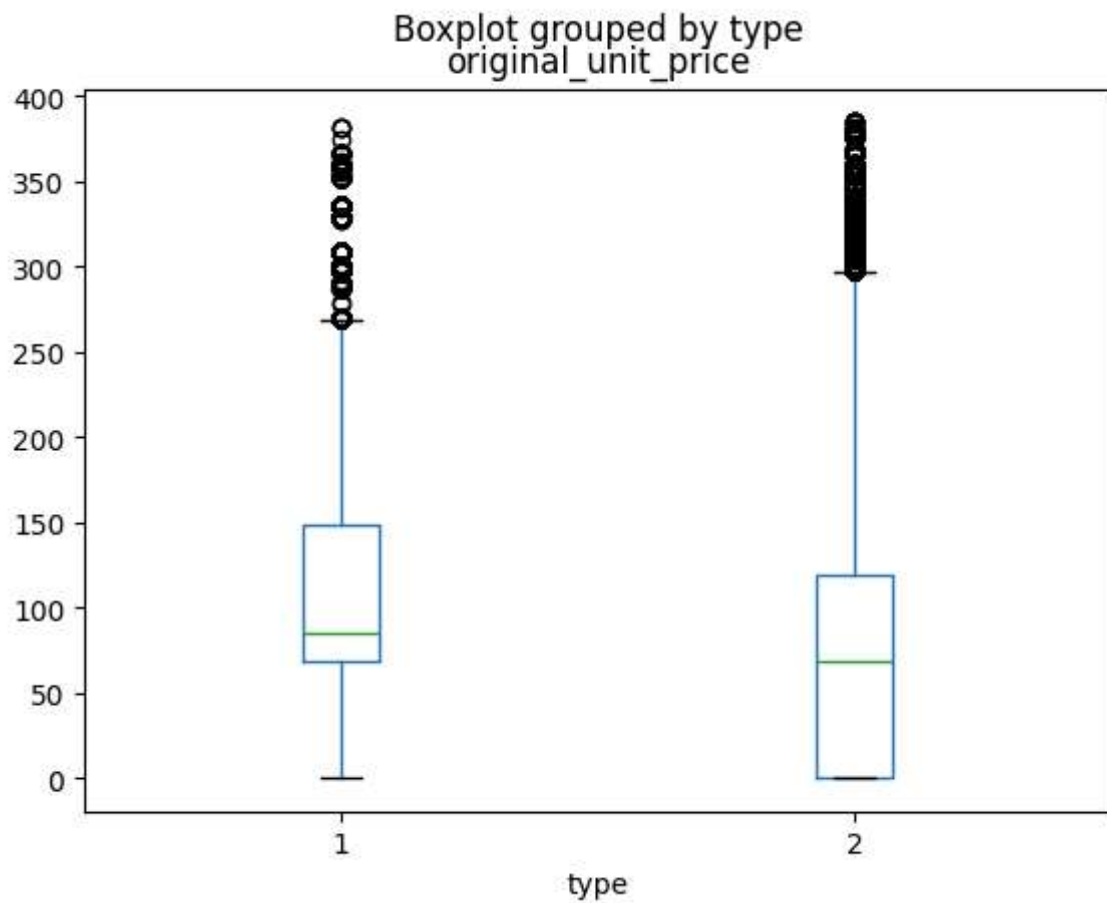
Out[52]: Text(0, 0.5, 'Final Price')



Exploring relationship between interval variable and categorical variable


```
In [53]: normal_price.boxplot(column = 'original_unit_price', by = 'type', grid = False)
         #price difference between products sold by jd and other third party
```

Out[53]: <Axes: title={'center': 'original_unit_price'}, xlabel='type'>



```
In [54]: sns.boxplot(x = 'type', y = 'original_unit_price', data = normal_price)
```

```
Out[54]: <Axes: xlabel='type', ylabel='original_unit_price'>
```



```
In [55]: sns.violinplot(x = 'type', y = 'original_unit_price', data = normal_price)
```

```
Out[55]: <Axes: xlabel='type', ylabel='original_unit_price'>
```

